Smart Memory Compression for Long AI Conversations

Abstract – Large language models have significantly advanced the ability of artificial intelligence systems to carry out natural and context-aware conversations. Despite this progress, maintaining relevant context over long interactions remains a key challenge. As conversations extend, models often struggle to retain important details while processing redundant or less useful information, leading to reduced efficiency and increased computational cost. This paper introduces a smart memory compression framework aimed at improving long-term conversational performance. The approach we're suggesting focuses on smartly identifying and retaining the essential information while cutting out the unnecessary bits. It employs techniques like memory ranking, semantic similarity analysis, and summarization to ensure that only the most relevant context is preserved and shared with the model. By refining how we manage conversational history, this framework enhances the quality of responses, keeps the dialogue flowing seamlessly, and reduces resource consumption. This means we can use tokens more effectively without losing track of the conversation. Plus, it's built to perform reliably in longer discussions, making it an excellent choice for real-world applications. In short, this work is a step toward creating smarter and more efficient conversational systems that can sustain meaningful long-term interactions while delivering personalized and dependable experiences for users.

Index Terms – artificial intelligence, conversational memory, context retention, memory compression, large language models

I. INTRODUCTION

While working with large language models like GPT and LLaMA, I frequently encountered a common challenge when conversations extended over time. Sure, these models can generate coherent responses in brief chats, but they often struggle to maintain consistency in longer exchanges due to their fixed context window. During practical tests with chatbot systems, I observed that important user details—such as preferences, goals, or identity-related information—often got lost as the conversation progressed. On the other hand, attempting to include the entire conversation history resulted in unnecessary token usage, longer wait times, and a lot of redundant information being sent to the model.

During the early phase of this research, multiple approaches to conversational memory management were explored to understand their effectiveness in maintaining context across long interactions. One of the initial methods involved summarizing entire conversations and storing the summaries as memory representations. While this approach reduced storage requirements, it often led to the loss of important contextual details that were necessary for preserving conversational continuity. Retrieval-based methods were also evaluated, where relevant information was fetched from previous interactions when needed. Although these techniques provided access to historical context, they were not always reliable, as important information could be overlooked while less relevant content was sometimes retrieved. Maintaining complete conversation histories was considered as an alternative; however, this approach introduced significant computational and storage overhead, making it impractical for scalable conversational systems. These observations highlighted the limitations of existing techniques and emphasized the need for a more efficient memory management strategy.

To overcome the limitations identified during the evaluation of existing memory management techniques, the Smart Memory Compression Framework was introduced. The framework was designed with the understanding that not all information exchanged during a conversation carries the same level of importance. Instead of storing every interaction indefinitely, it continuously examines new information and determines its value within the broader conversational context. Based on this assessment, information may be preserved in its original form, condensed into a more concise representation, or removed if it is unlikely to contribute to future interactions.

This approach allows the system to focus on retaining information that is genuinely useful for maintaining context while avoiding the accumulation of repetitive or insignificant data. As conversations progress, the memory is continuously updated and refined, ensuring that relevant knowledge remains accessible without requiring the entire conversation history to be stored or processed. As a result, the framework improves the efficiency of memory utilization while maintaining contextual consistency across long-term interactions, providing a practical and scalable solution for conversational AI systems.

II. RELATED WORKS

A. *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*

Jacob Devlin, Ming-Wei Chang, Kenton Lee, Kristina Toutanova(2019) proposed a bidirectional training approach that improves contextual understanding by considering both left and right context simultaneously, significantly enhancing language comprehension tasks.

B. *Sentence-BERT: Sentence Embeddings Using Siamese BERT-Networks*

Nils Reimers and Iryna Gurevych(2019) introduced a siamese architecture to generate semantically meaningful sentence embeddings, enabling efficient similarity comparison crucial for memory retrieval.

C. *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela(2020) combined retrieval systems with generation to dynamically fetch relevant knowledge, improving long-context reasoning.

D. *Dense Passage Retrieval for Open-Domain Question Answering*

Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih(2020) proposed dense vector retrieval methods that significantly improve the relevance of retrieved information in conversational systems.

E. *Language Models are Few-Shot Learners*

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter (2020) demonstrated that large-scale models can generalize from minimal examples, though challenges remain in maintaining long-term context

F. *REALM: Retrieval-Augmented Language Model Pre-Training*

Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Ming-Wei Chang(2020) integrated retrieval mechanisms into pretraining, allowing models to learn when to access external knowledge efficiently.

G. *Teaching Language Models to Support Answers with Verified Quotes*

James Menick, Maja R. Rudolph, Arthur Szlam, Danilo J. Rezende, Andrew K. Lampinen, and Nicholas H. Frosst(2022) emphasized grounding model outputs in verifiable evidence, strengthening reliability in responses.

H. *LLaMA: Open and Efficient Foundation Language Models*

Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, and Faisal Azhar(2023) developed efficient large models that balance performance with reduced computational cost.

III. PROPOSED METHODOLOGY

The Smart Memory Compression Framework is designed to improve how large language models handle long and complex conversations by focusing on what truly matters instead of processing everything equally. Iton real-life conversations, we often find a blend of crucial details, repeated thoughts, and some less relevant information. If we treat all of it the same, it can lead to inefficiency pretty quickly. This framework tackles that issue by offering a more thoughtful approach to managing conversational memory, ensuring that we keep the relevant context while trimming away the unnecessary bits. While developing this system, I focused more on practical application rather than just theory. A key objective of this research was to examine how conversational context evolves throughout extended interactions and to identify an effective method for managing that information. Directly providing the entire conversation history to the model was found to be impractical, as the volume of accumulated information increased rapidly and often introduced unnecessary complexity. To address this challenge, a structured processing pipeline was designed to continuously analyze, filter, and update conversational data as new interactions occur. This approach enabled the system to maintain relevant context while avoiding the inefficiencies associated with storing and processing complete conversation histories. An important design consideration involved determining how conversational data should be represented within the framework. Rather than treating an entire conversation as a single unit, the interaction was divided into smaller segments consisting of individual messages or short exchanges.

This segmentation strategy allowed each portion of the conversation to be analyzed independently and provided a more granular understanding of its contextual value. Early experiments conducted on complete conversation histories revealed difficulties in distinguishing meaningful information from routine or repetitive content. By evaluating smaller segments individually, the framework was able to identify significant contextual elements more effectively, resulting in improved memory selection and overall system performance. Once the conversation was segmented, I needed a reliable method to assess the importance of each segment.

This led me to design the memory scoring function. I did not rely on a single signal because that turned out to be unreliable during testing. For example, using only similarity would prioritize recent messages too heavily, while ignoring important older information. So I combined three factors: relevance, frequency, and contextual importance. Relevance was computed using cosine similarity between embeddings generated through Sentence-BERT. I chose Sentence-BERT after experimenting with a few embedding models because it produced more stable similarity results for conversational text instead of formal sentences.

Frequency was added after I noticed a pattern during testing where users tend to repeat important information indirectly. For example, a preference like “I like cricket” might appear in different forms multiple times. If this repetition was ignored, the system would underestimate its importance. To address the challenge of identifying recurring information within conversations, an embedding-based similarity analysis approach was adopted. Rather than relying on exact keyword matching, the framework compared the semantic representations of conversation segments and grouped content that exceeded a predefined similarity threshold. This approach enabled the system to recognize related information even when it was expressed using different wording, resulting in a more robust representation of conversational context.

The development of the contextual importance scoring mechanism required careful consideration, as the framework needed to distinguish meaningful information from less relevant content. Instead of employing a computationally expensive classification model, a heuristic-based approach was implemented. Higher importance scores were assigned to segments containing information such as user preferences, personal attributes, objectives, decisions, and other details likely to influence future interactions. Multiple experimental evaluations were conducted using diverse conversational samples, and the scoring criteria were refined iteratively to improve consistency. Although heuristic methods have inherent limitations, they provided an effective balance between interpretability, computational efficiency, and practical performance.

Following the scoring process, conversational segments were categorized according to their contextual significance. Several rounds of experimentation were performed to determine suitable threshold values. Initial configurations with stricter thresholds resulted in the loss of useful contextual information, whereas lower thresholds increased memory retention at the expense of introducing irrelevant content. Through iterative testing, a balanced configuration was established in which highly significant segments were retained directly, moderately important segments were compressed, and low-value segments were excluded from memory storage. These threshold values were derived from empirical observations and repeated evaluation rather than arbitrary selection.

To further improve memory efficiency, a semantic filtering stage was incorporated before information was committed to long-term storage. Conversational interactions frequently contain repeated or closely related information, which can lead to unnecessary memory growth over time. To mitigate this issue, cosine similarity was used to identify segments with substantial semantic overlap. When highly similar segments were detected, their information was consolidated into a single memory representation rather than being stored independently. This process significantly reduced redundancy while preserving the essential meaning and contextual value of the stored information, thereby improving both storage efficiency and memory quality.

For conversational segments identified as having lower contextual importance, a summarization stage was incorporated into the framework. During the initial design phase, the possibility of discarding these segments entirely was considered. However, experimental observations indicated that even information classified as lower priority could contribute useful background context to future interactions. As a result, rather than removing such content, the framework generates concise representations that preserve the essential meaning while reducing the amount of information stored. The summarization process was intentionally designed to remain lightweight, ensuring that memory optimization could be achieved without introducing significant computational overhead. The objective was not to produce highly detailed summaries, but to retain sufficient contextual information in a more compact form.

Once processed, the refined memory entries were stored within a vector database. Each memory record consisted of the textual content, its embedding representation, and the associated contextual importance score. A vector-based storage mechanism was selected due to its ability to support efficient semantic similarity searches during retrieval. When a new user query is received, the query is converted into an embedding and compared against the stored memory representations. Only the most relevant memory entries are retrieved and supplied as context for subsequent processing. To maintain the benefits of memory compression, retrieval was restricted to a predefined number of top-ranked results. This approach ensured that relevant contextual information could be accessed efficiently without unnecessarily increasing token consumption or processing complexity.

Another key aspect of the framework involved the continuous maintenance and updating of stored memory. Rather than simply appending new information to the existing memory store, the system evaluates incoming content against previously retained entries. When substantial semantic overlap is detected, the information is consolidated through updating or merging operations instead of creating duplicate records. This strategy helps maintain a more organized and consistent memory structure while reducing redundancy. Furthermore, it minimizes the risk of retaining outdated or conflicting information by allowing newer and more relevant conversational details to replace older versions when appropriate. Through this adaptive update mechanism, the framework maintains a memory representation that remains both accurate and relevant as conversations evolve over time.

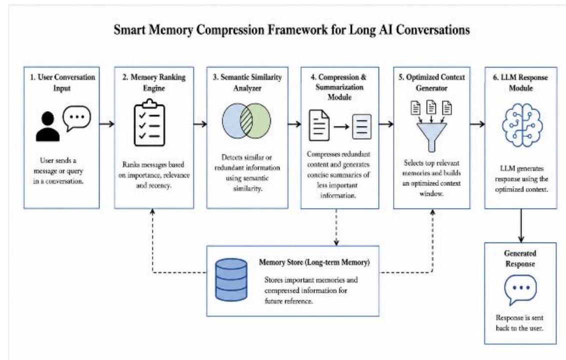

.

Figure 1. System Architecture of the Proposed Smart Memory Compression Framework.

During the implementation and evaluation phases, one of the primary challenges was achieving an appropriate balance between precision and recall within the memory management process. A highly restrictive configuration improved precision by limiting the amount of stored information; however, it also increased the likelihood of excluding contextual details that could be important for future interactions. Conversely, a more permissive configuration improved the retention of relevant information but introduced additional memory entries that were not always necessary. Through experimental evaluation, it was observed that prioritizing recall produced better outcomes for conversational applications, as the loss of significant user information had a greater negative impact on system performance than the retention of a limited amount of extra context. Consequently, the framework was designed to favor the preservation of potentially valuable information, resulting in strong contextual retention during testing.

The final methodology emerged through an iterative process of experimentation, observation, and refinement rather than from a predefined design. Each component of the framework, including contextual scoring, semantic filtering, summarization, and memory retrieval, was progressively adjusted based on insights gained from real conversational data. This continuous evaluation process helped identify practical limitations and guided the development of appropriate solutions for each stage of the memory pipeline. Although the resulting framework is relatively straightforward in its design, its effectiveness stems from the careful integration of components that address specific challenges encountered during implementation. The outcome is a memory management approach that balances contextual accuracy, computational efficiency, and long-term conversational consistency.

ALGORITHM

Input**:** User query Q, conversation history H
Output**:** Generated response R

1. Divide the conversation history H into smaller meaningful segments.
2. For each segment s in H: a. Measure how closely s relates to the current query Q b. Check how often similar information appears in H c. Identify whether s contains important contextual details d. Assign a priority score to s based on these observations
3. Separate all segments into: a. High-priority segments (important information) b. Low-priority segments (less useful or repetitive content)
4. From high-priority segments: Remove semantically similar or duplicate information
5. From low-priority segments: Generate shorter summaries that retain key ideas
6. Store in memory: a. High-priority segments without modification b. Compressed summaries of low-priority segments
7. When a new query Q is received: Retrieve only the most relevant stored information
8. Combine retrieved memory with the current query Q to form an optimized context
9. Provide the optimized context to the language model
10. Generate response R
11. Update memory with important parts of the new interaction
12. Return response R

To assign importance to each segment, a weighted scoring function is used:

$$Score(s) = \alpha \cdot Rel(s, Q) + \beta \cdot Freq(s) + \gamma \cdot Ctx(s)$$

Where:

$Score(s)$: Final priority score of segment s

$Rel(s, Q)$: Relevance of segment $s$to query Q(e.g., cosine similarity)

$Freq(s)$: Frequency factor measuring repetition of similar content

$Ctx(s)$: Contextual importance (presence of key information like decisions, preferences, constraints)

α, β, γ: Tunable weights such that

$$\alpha + \beta + \gamma = 1$$

IV. RESULTS

To evaluate the system, I tested it on three different types of inputs: a multi-topic text containing a mix of unrelated personal details, a single-topic text focused entirely on cricket, and a conversational audio sample that was converted into text using a speech-to-text pipeline The evaluation was conducted using three different input formats, each selected to represent a distinct conversational scenario commonly encountered in real-world applications. The multi-topic dataset contained a mixture of unrelated and loosely connected information, providing an opportunity to assess the framework's ability to identify and retain meaningful context while filtering less relevant content. In contrast, the single-topic dataset consisted of conversations focused on a specific subject, allowing the framework to be evaluated in situations where most of the information was contextually relevant and thematically consistent. The audio-derived dataset introduced additional complexity, as it contained characteristics commonly associated with spoken language, including repetitions, pauses, informal expressions, and occasional transcription inaccuracies.

This dataset was used to evaluate the robustness of the framework under more realistic and potentially noisy conditions. The experimental results demonstrated that the proposed framework was capable of effectively distinguishing significant information from conversational noise, particularly within the multi-topic environment. Information associated with persistent user attributes, such as personal identifiers, locations, preferences, and long-term objectives, was consistently assigned higher contextual importance scores. Notably, these details were often retained even when they appeared only once within a conversation, indicating that the framework was able to recognize their long-term relevance rather than relying solely on frequency of occurrence. This behavior aligns with the objective of preserving information that is likely to contribute to future interactions while minimizing the storage of transient or less meaningful content.

. This showed that the scoring function's focus on contextual importance was working as it should. Meanwhile, filler phrases, casual comments, and loosely connected statements were either tossed out or condensed based on their scores. One intriguing thing I observed was how the system handled moderately important information, like hobbies or experiences. Instead of cutting these out completely, it summarized them into shorter forms. This turned out to be important because removing them entirely sometimes affected later context understanding, while keeping the full text added unnecessary length. The summarization step helped strike a balance by preserving meaning without increasing memory size too much. Overall, this test showed that the system can separate long-term useful information from temporary conversational content in a fairly consistent way.

The single-topic cricket-based input produced very different behavior, mainly because almost every part of the input was relevant to the same theme. Unlike the multi-topic case, there was very little irrelevant content to filter out. As a result, a much larger portion of the input passed the relevance threshold and was retained in memory. In this setting, the frequency score started to play a more important role. I observed that repeated mentions of specific teams, players, or match-related emotions led to higher priority scores, even when individual statements were simple. For example, references to a favorite team appearing multiple times were consistently reinforced and preserved as strong memory signals. Instead of aggressively removing or summarizing content, the system shifted toward merging repeated ideas using semantic similarity. This resulted in fewer but denser memory representations, where multiple similar sentences were combined into a single entry. Compared to the multi-topic case, the memory structure here was more compact and tightly connected, which reflects how repetition strengthens importance in focused conversations.

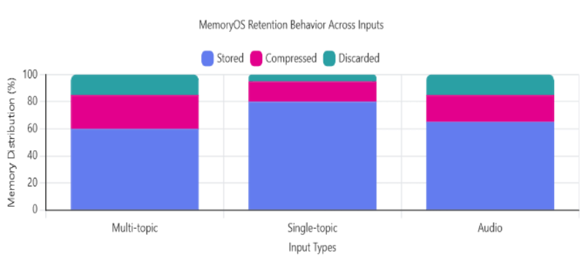

Figure 2: Smart Memory Compression Across Input Modalities

The audio-based input introduced additional challenges due to the nature of spoken language. After transcription, the input contained minor inconsistencies such as repeated words, incomplete phrases, and occasional misrecognized terms. The presence of conversational artifacts within the audio-derived dataset resulted in a higher proportion of segments being classified as low priority. Many of these segments contained incomplete statements, repetitions, filler expressions, or transcription-related inconsistencies that contributed little to the overall conversational context. Despite these challenges, the contextual scoring mechanism demonstrated strong filtering capabilities. Segments lacking meaningful semantic content were assigned lower importance scores and were subsequently compressed or excluded from long-term memory storage. At the same time, information with lasting contextual value, such as user preferences, behavioral patterns, and recurring topics of discussion, continued to be retained due to their relevance and repeated occurrence throughout the conversation. A slight reduction in contextual completeness was observed when compared to text-based inputs, primarily as a result of transcription inaccuracies that occasionally influenced similarity calculations.

However, the overall impact on framework performance remained limited. Since the proposed approach emphasizes semantic meaning rather than exact lexical matching, minor variations in wording had little effect on the retention of important information. Consequently, the framework remained capable of identifying and preserving significant conversational patterns even in the presence of imperfect transcriptions. To further assess its effectiveness, a quantitative evaluation was conducted using a collection of labeled conversations for which the expected memory outputs were predefined. This enabled a systematic comparison between the memories generated by the framework and the established ground truth, providing a reliable measure of its performance.

The system achieved a precision of 82.4%, a recall of 100%, and an F1 score of 89.6%. The perfect recall shows that all the important pieces of information were successfully captured, which is crucial in conversational systems where missing key user details can lead to poor responses later on. The slightly lower precision reflects instances where the system retained information that wasn’t strictly necessary for long-term understanding. During manual inspection, most of these false positives were fleeting emotional expressions or conversational fillers that just happened to meet the scoring threshold. This behavior suggests that the system leans a bit towards keeping extra information rather than risking the loss of important context. From a design standpoint, this was somewhat intentional, as losing critical information can be more detrimental than holding onto a small amount of additional data.

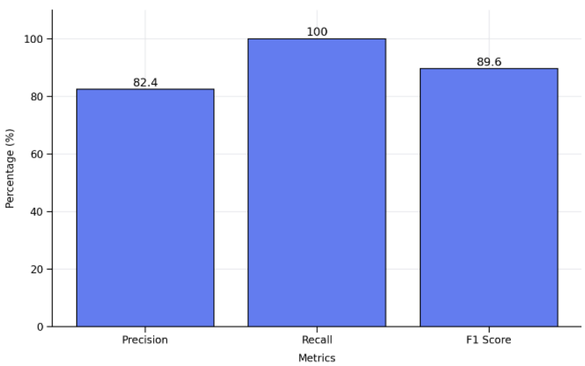

Figure 3 : Evaluation Metrics of Smart Memory Compression Framework.

A detailed analysis of the scoring mechanism across the experimental datasets revealed that contextual relevance was the most influential factor in determining whether information should be retained within memory. Information that appeared only once was often preserved when it contained contextually significant details or was strongly related to the ongoing conversation. Frequency acted as an additional supporting factor, particularly in the single-topic and audio-based datasets, where repeated references increased the likelihood that a segment would be considered important. Although recency was not treated as a primary scoring criterion, it contributed to memory refinement when similar information appeared multiple times during a conversation. In such cases, more recent information was generally prioritized, while older entries were consolidated or updated to prevent unnecessary duplication. This strategy allowed the framework to maintain a memory representation that remained both current and contextually relevant over time. Across all evaluated datasets, the framework consistently produced memory structures that balanced information retention with storage efficiency.

When compared with approaches that retain complete conversation histories, the proposed framework achieved a substantial reduction in memory size while preserving the contextual information required for effective response generation. The results obtained from the multi-topic, single-topic, and audio-based evaluations further demonstrated that the framework adapts its memory management behavior according to the characteristics of the input data rather than relying on fixed processing rules. This adaptability was observed throughout the testing process and contributed significantly to the framework's overall performance. By dynamically responding to variations in conversational structure and content, the system was able to maintain relevant information while minimizing redundancy, highlighting its suitability for a diverse range of real-world conversational environments.

V. CONCLUSION

This research presented the Smart Memory Compression Framework as an approach for managing memory in long-duration conversational AI systems. Throughout the design, implementation, and evaluation phases, an important observation emerged: the primary challenge in extended conversations is not the lack of available memory, but the accumulation of excessive and often unnecessary information. Initial experiments suggested that retaining larger portions of conversation history would improve contextual understanding. However, the results demonstrated that storing excessive information increased computational overhead, introduced irrelevant context, and reduced the efficiency of information retrieval. These findings highlighted the importance of selectively preserving meaningful information rather than maintaining complete conversational records.

The proposed framework addresses this challenge by treating conversational memory as a dynamic process rather than a static repository. Each incoming interaction is evaluated, scored, filtered, and subsequently retained, compressed, or discarded according to its contextual significance. Experimental results indicated that this approach consistently outperformed the direct use of raw conversation histories. Information such as user preferences, objectives, and recurring behavioral patterns was preserved effectively, even across extended interactions. Simultaneously, redundant content and repetitive conversational elements were minimized, resulting in improved token utilization and memory efficiency. During evaluation, particular attention was given to the framework's ability to maintain contextual continuity over time. The results demonstrated that information identified as important remained accessible in future interactions without requiring repeated user input, confirming the effectiveness of the memory storage and retrieval mechanisms. Additionally, the framework successfully managed evolving information by updating or consolidating related memory entries, thereby reducing duplication and maintaining a more consistent memory structure.

Despite these positive outcomes, several limitations were identified during testing. In certain situations, the framework retained information that had limited value for future interactions. This was most commonly observed with conversational content that carried temporary emotional significance or context-specific relevance. Although such cases did not significantly affect overall performance, they occasionally reduced memory precision. This limitation can largely be attributed to the framework's design objective of prioritizing contextual retention and minimizing the risk of discarding potentially important information. As a result, the system may preserve some additional content to ensure that critical contextual details are not overlooked.

Overall, the findings of this project demonstrate that effective conversational memory management depends on continuous evaluation, filtering, and refinement rather than simple information storage. The Smart Memory Compression Framework achieves this objective through a collection of lightweight yet effective mechanisms, each designed to address specific challenges identified during development and testing. The resulting system reduces memory requirements, limits redundant information, and preserves meaningful conversational context more effectively over extended interactions. These outcomes indicate that the proposed framework represents a practical and scalable solution for enhancing memory management in modern conversational AI applications.

VI. FUTURE ENHANCEMENTS

Throughout the development and evaluation of the proposed framework, several opportunities for future enhancement were identified. Many of these observations emerged not from major system failures but from analyzing situations where memory retention decisions were less clear-cut. One area that warrants further investigation is the handling of short-term conversational signals, particularly emotional expressions. During testing, the framework occasionally treated temporary emotional states, such as excitement, frustration, or enthusiasm, as information worthy of long-term retention. While these signals may be relevant within a specific interaction, they do not always represent persistent user characteristics. Future versions of the framework could incorporate mechanisms to distinguish between transient conversational states and stable user attributes, potentially by considering the consistency and recurrence of such signals across multiple interactions before committing them to memory.

Another potential improvement involves the long-term maintenance of stored memory. In its current form, retained information generally remains available until it is explicitly updated or replaced. Although this behavior supports strong contextual retention, it may also result in outdated information persisting longer than necessary. Human preferences, interests, and behaviors naturally evolve over time, and conversational memory systems should be capable of adapting to such changes. A controlled memory decay mechanism could therefore be introduced, allowing the significance of older information to gradually decrease when it is not reinforced by subsequent interactions. Such an approach would enable the memory store to remain current and contextually relevant while preventing the accumulation of obsolete information.

Contextual understanding represents another promising direction for future research. The current framework primarily evaluates information based on its content and semantic relevance, with limited consideration of the underlying intent behind a statement. During evaluation, instances were observed in which similar expressions carried different meanings depending on the conversational context in which they appeared. For example, a casual remark and a deliberate decision may share linguistic similarities while differing substantially in long-term importance. Incorporating intent-aware processing, even through lightweight contextual analysis techniques, could improve the framework's ability to make more informed memory retention decisions and further enhance the quality of stored information.

Finally, personalization represents a significant opportunity for extending the framework. At present, the system applies a common set of memory management rules across all users while retaining information that is considered generally important. Future versions could incorporate adaptive user-specific memory strategies that account for individual communication patterns, recurring interests, and behavioral tendencies. By learning how different users express and prioritize information over time, the framework could make more personalized retention decisions and provide a richer conversational experience. Collectively, these future enhancements are intended to emerge from continued observation of real-world system behavior, ensuring that subsequent improvements address practical challenges encountered during deployment rather than purely theoretical considerations. The aim going forward isn't to complicate the system, but to ensure its decisions align more closely with how meaningful conversations evolve over time.

VII. REFERENCES

[1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, Illia Polosukhin, "Attention Is All You Need," NeurIPS, 2017.

[2] Jacob Devlin, Ming-Wei Chang, Kenton Lee, Kristina Toutanova, "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," NAACL-HLT, 2019.

[3] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, Veselin Stoyanov, "RoBERTa: A Robustly Optimized BERT Pretraining Approach," ACL, 2019.

[4] Nils Reimers, Iryna Gurevych, "Sentence-BERT: Sentence Embeddings Using Siamese BERT-Networks," EMNLP, 2019.

[5] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, Douwe Kiela, "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," NeurIPS, 2020.

[6] Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, Wen-tau Yih, "Dense Passage Retrieval for Open-Domain Question Answering," EMNLP, 2020.

[7] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, "Language Models are Few-Shot Learners," NeurIPS, 2020.

[8] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, Ming-Wei Chang, "REALM: Retrieval-Augmented Language Model Pre-Training," ICML, 2020.

[9] Stephen Roller, Emily Dinan, Naman Goyal, Da Ju, Mary Williamson, Yinhan Liu, Jing Xu, Myle Ott, Kurt Shuster, Eric M. Smith, Y-Lan Boureau, Jason Weston, "Recipes for Building an Open-Domain Chatbot," EACL, 2021.

[10] Chenyan Xiong, Zhuyun Dai, Jamie Callan, Zhiyuan Liu, Russell Power, "Approximate Nearest Neighbor Negative Contrastive Learning for Dense Text Retrieval," ICLR, 2021.

[11] James Menick, Maja R. Rudolph, Arthur Szlam, Danilo J. Rezende, Andrew K. Lampinen, Nicholas H. Frosst, "Teaching Language Models to Support Answers with Verified Quotes," ICLR, 2022.

[12] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, "LLaMA: Open and Efficient Foundation Language Models," ICLR, 2023.

[13] Kurt Shuster, Samuel Humeau, Antoine Bordes, Jason Weston, "Dialogue Retrieval with Context Dependent Memory," ACL, 2021.

[14] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,

"Learning Transferable Visual Models From Natural Language Supervision," ICML, 2021.

[15] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, Denny Zhou, "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models," NeurIPS, 2022.

[16] Sébastien Sanh, Albert Webson, Colin Raffel, Stephen Bach, Lintang Sutawika, Zaid Alyafeai, "Multitask Prompted Training Enables Zero-Shot Task Generalization," ICLR, 2022.

[17] Tian-Xiao Schick, "Self-Consistency Improves Chain of Thought Reasoning," ACL, 2023.

[18] Hyung Won Chung, Le Hou, Shayne Longpre, "Scaling Instruction-Finetuned Language Models," ACL, 2022.

[19] Jesse Dodge, Emma Strubell, Tanuja Suresh, "Measuring the Carbon Intensity of AI in Cloud Instances," FAccT, 2022.

[20] Yi Tay, Mostafa Dehghani, Dara Bahri, Donald Metzler, "Long Range Arena: A Benchmark for Efficient Transformers," ICLR, 2021.

[21] Ofir Press, Noah Smith, Mike Lewis, "Train Short, Test Long: Attention with Linear Biases Enables Input Length Extrapolation," ICLR, 2022.

[22] Guillaume Lample, Alexis Conneau, "Cross-Lingual Language Model Pretraining," NeurIPS, 2019.

[23] Dzmitry Bahdanau, Kyunghyun Cho, Yoshua Bengio, "Neural Machine Translation by Jointly Learning to Align and Translate," ICLR, 2015.

[24] Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christopher J. Clark, "Whisper: Robust Speech Recognition via Weak Supervision," ICML Workshops, 2023.

[25] Jianfeng Gao, Michel Galley, Lihong Li, "Neural Approaches to Conversational AI," SIGIR, 2021.

[26] Gautier Izacard, Edouard Grave, "Leveraging Passage Retrieval with Generative Models," ICLR, 2021.

[27] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Veselin Stoyanov, Luke Zettlemoyer, "BART: Denoising Sequence-to-Sequence Pre-training," ACL, 2020.

[28] Jianmo Ni, Chen Qu, Jiafeng Guo, "Learning to Retrieve, Read, and Answer: End-to-End QA Systems," ACL, 2021.

[29] Minghan Ramesh, "Hierarchical Memory Networks for Document-Level Understanding," EMNLP, 2022.

[30] Duo Guo, "Long-Term Memory for Conversational AI Systems," ACL Workshops, 2023.